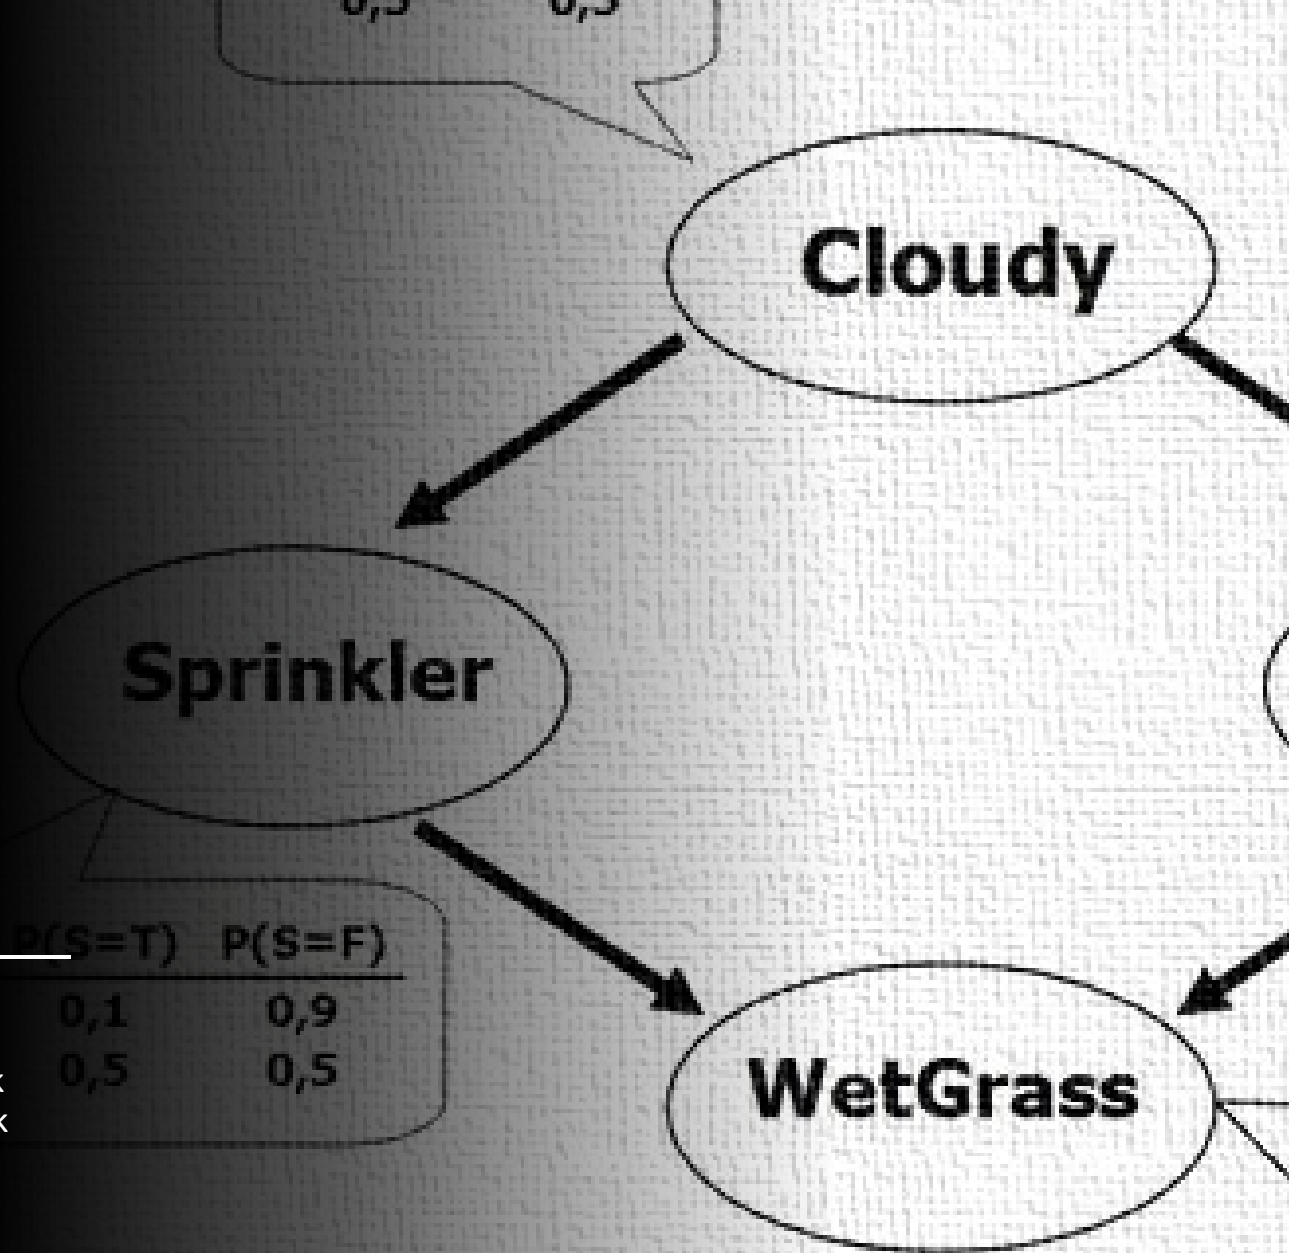


CS 5/7320 Artificial Intelligence

Probabilistic Reasoning AIMA Chapter 13

Slides by Michael Hahsler
based on slides by Svetlana Lazepnik
with figures from the AIMA textbook



Probability Recap

- **Notation:** $P(X)$ vs. $P(X = x) = P(x)$

- Conditional probability $P(x|y) = \frac{P(x, y)}{P(y)} = \alpha P(x, y)$

- Product rule $P(x, y) = P(x|y)P(y)$


- Chain rule $P(X_1, X_2, \dots, X_n) = P(X_1)P(X_2|X_1)P(X_3|X_1, X_2) \dots$
 $= \prod_{i=1}^n P(X_i|X_1, \dots, X_{i-1})$

- X, Y are independent if and only if: $\forall x, y : P(x, y) = P(x)P(y)$
Written as $X \perp\!\!\!\perp Y$

- X and Y are conditionally independent given Z if and only if:

$$\forall x, y, z : P(x, y|z) = P(x|z)P(y|z) \quad \text{Written as } X \perp\!\!\!\perp Y|Z$$

Contents



Cloudy

Sprinkler

Rain

WetGrass

Bayesian Networks

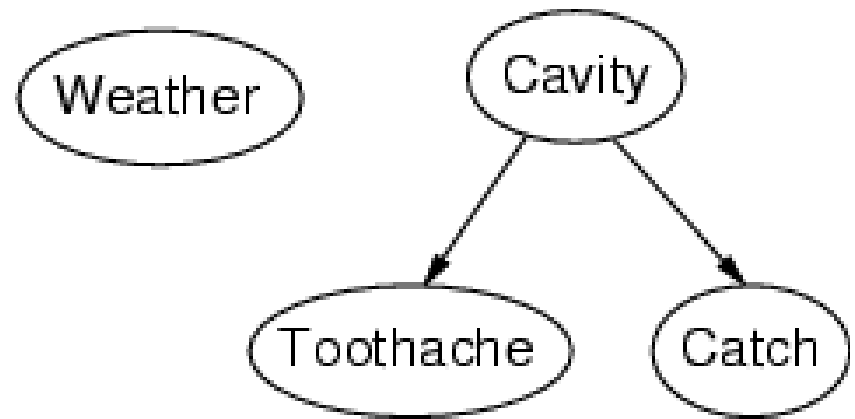
Modeling a Joint Distribution

C	P
T	
F	

C	P(S=T)	P(S=F)
T	0,1	0,9
F	0,5	0,5

S	R	P(W)
T	T	0,9
T	F	0,5
F	T	0,5
F	F	0,1

Bayesian Networks (aka Belief Networks)



A type of graphical model.



A way to specify dependence between random variables.



A compact specification of a full joint distributions.

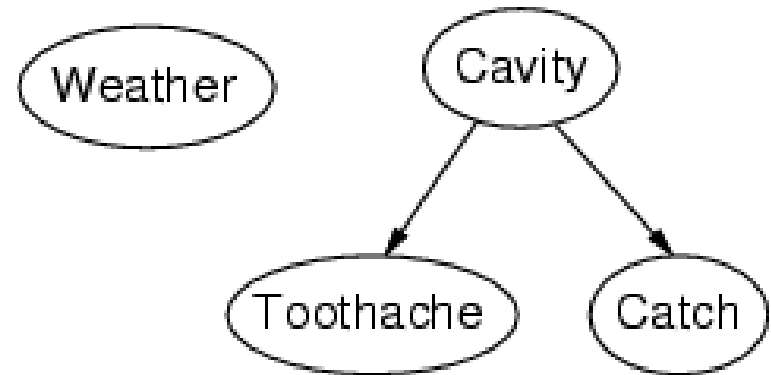


A general and important model to reason with uncertainty in AI.

Structure of Bayesian Networks

Nodes: Random variables

- Can be assigned (observed) or unassigned (unobserved)




Arcs: Dependencies

- An arrow from one variable to another indicates direct influence.
- Show independence
 - *Weather* is independent of the other variables (no connection).
 - *Toothache* and *Catch* are conditionally independent given *Cavity* (directed arc).
- Must form a directed *acyclic* graph (DAG)

A network with all random variables assigned represents a **state of the system**.

Example: N independent coin flips

Complete independence: no interactions between coin flips



$$P(X_1, X_2, \dots, X_n) = P(X_1)P(X_2) \dots P(X_n)$$

Joint probability
distribution

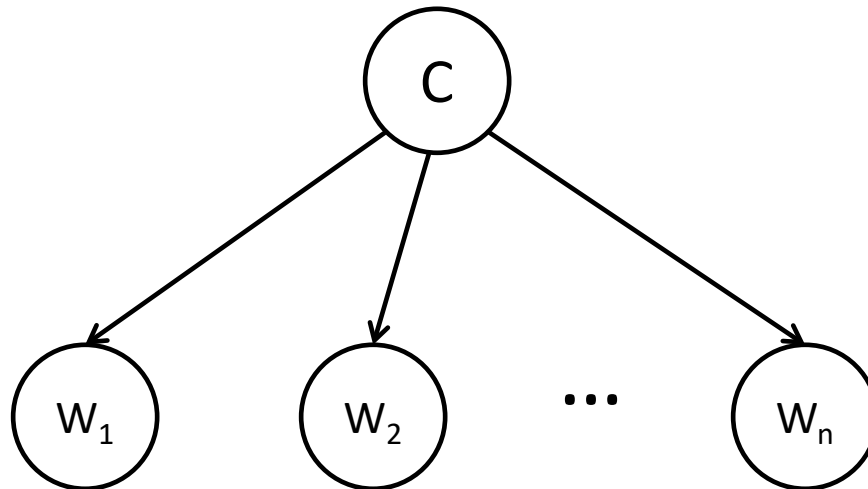
Marginal probability
distributions

Example: Naïve Bayes spam filter

Random variables:

- C : message class (spam or not spam)
- W_1, \dots, W_n : presence or absence of words comprising the message

Words depend on the class, but they are modeled conditional independent of each other given the class (= no direct connection between words).




$$P(W_1, W_2, \dots, W_n | C) = P(W_1 | C) P(W_2 | C) \dots P(W_n | C)$$

Example: Burglar Alarm

- **Description:** I have a burglar alarm that is sometimes set off by minor earthquakes. My two neighbors, John and Mary, promised to call me at work if they hear the alarm
- Example inference task: suppose Mary calls and John doesn't call. What is the probability of a burglary?
- What are the random variables?
 - Burglary, Earthquake, Alarm, JohnCalls, MaryCalls
- What are the direct influence relationships?
 - A burglar can set off the alarm
 - An earthquake can set off the alarm
 - The alarm can cause Mary to call
 - The alarm can cause John to call


Example: Burglar Alarm



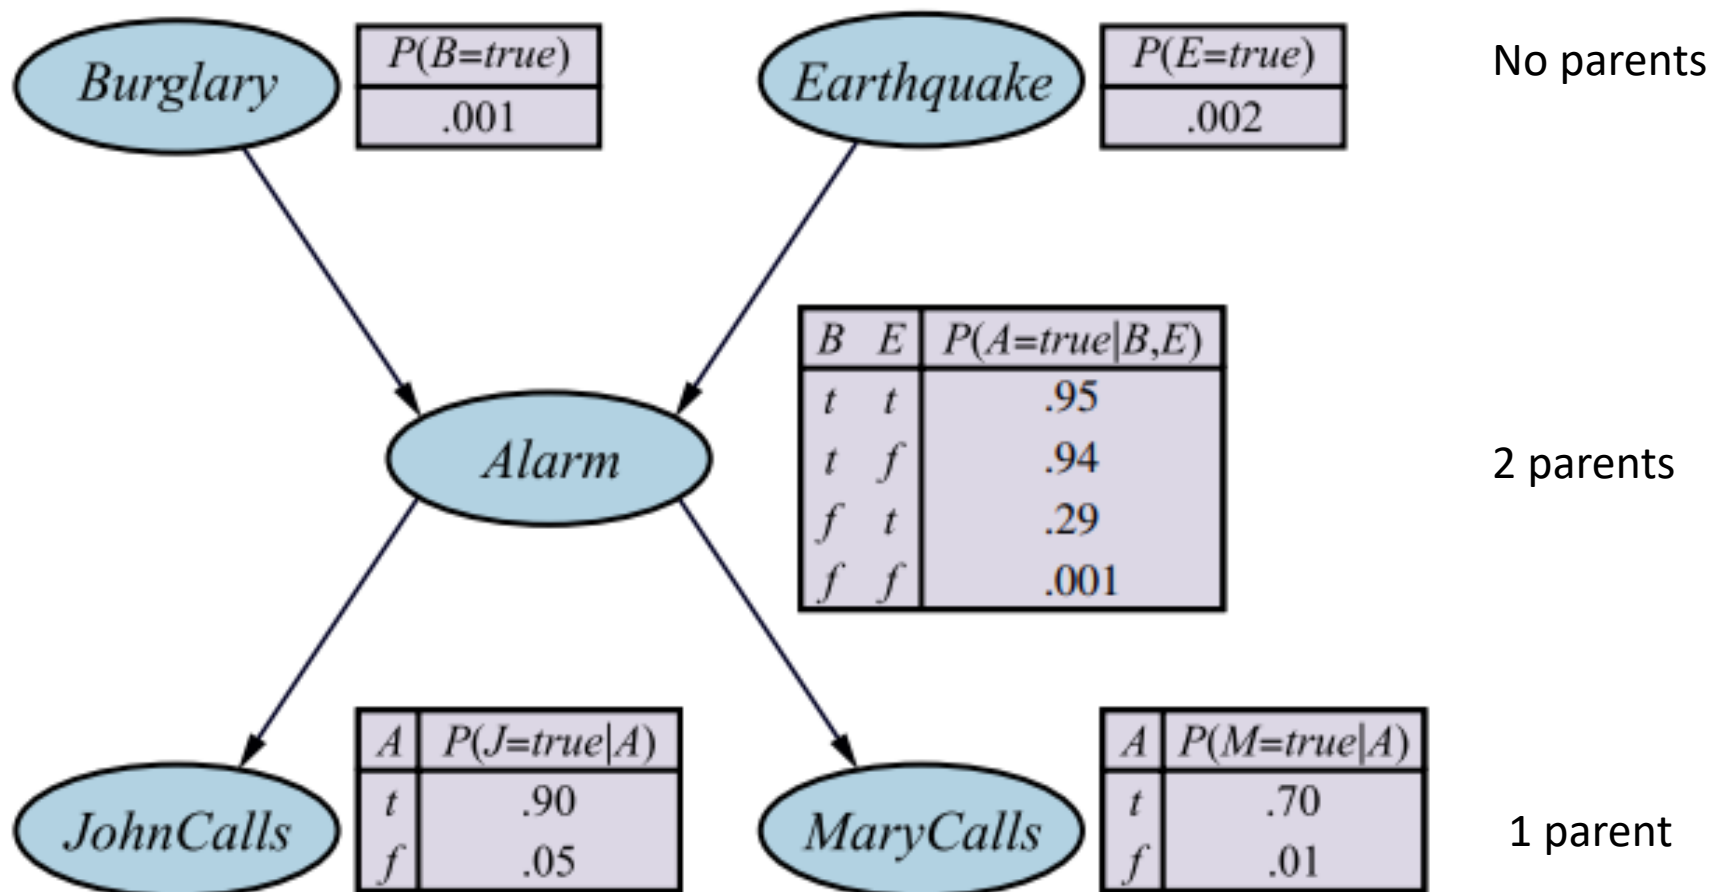
What are the model parameters?

Parameters: Conditional probability tables

To specify the full joint distribution, we need to specify a *conditional* distribution for each node given its parents as a conditional probability table (CPT): $P(X | \text{Parents}(X))$



Example: Burglar Alarm with CPTs



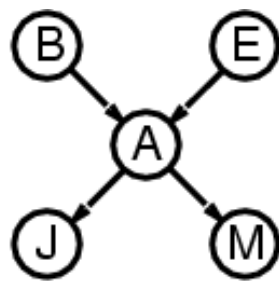
The joint probability distribution

- For each node X_i , we know $P(X_i \mid \text{Parents}(X_i))$
- How do we get the full joint distribution $P(X_1, \dots, X_n)$?

- Using chain rule:

$$P(X_1, \dots, X_n) = \prod_{i=1}^n P(X_i \mid X_1, \dots, X_{i-1}) = \prod_{i=1}^n P(X_i \mid \text{Parents}(X_i))$$

- Example:

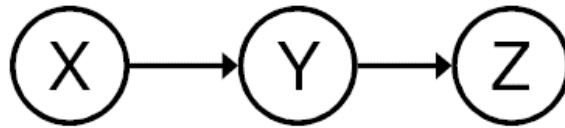


Construct
following
arrows

$$P(J, M, A, B, E) = P(B) P(E) P(A \mid B, E) P(J \mid A) P(M \mid A)$$

Dependence

- Example: *causal chain*



X: Low pressure

Y: Rain

Z: Traffic

- Are X and Z independent?

$$P(X, Y, Z) = P(X)P(Y|X)P(Z|Y)$$

Conditioning

$$P(X, Z) = \sum_y P(X)P(y|X)P(Z|y)$$

Marginalize
over Y

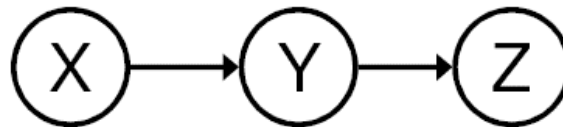
$$= P(X) \sum_y P(Z|y)P(y|X) \neq P(X)P(Z)$$



X and Z are not independent!

Conditional independence

- Example: *causal chain*



X: Low pressure

Y: Rain

Z: Traffic

- Is Z independent of X given Y?

$$P(X, Z|Y) = \frac{P(X, Y, Z)}{P(Y)} = \frac{P(X)P(Y|X)P(Z|Y)}{P(Y)}$$

Conditioning

$$= \frac{P(X) \frac{P(X|Y)P(Y)}{P(X)} P(Z|Y)}{P(Y)}$$

Bayes' rule

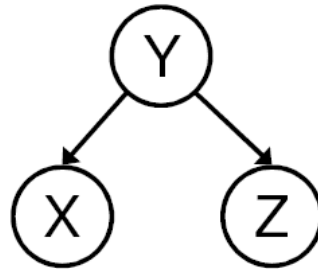
$$= P(X|Y)P(Z|Y) = \text{Definition of conditional independence}$$



X and Z are conditionally independent given Y

Conditional independence

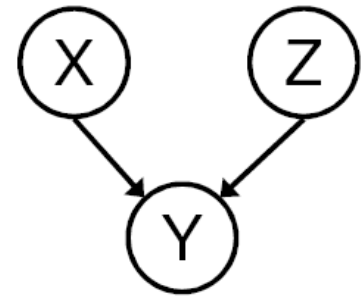
- Common cause



Y: Project due
X: Newsgroup busy
Z: Lab full


- Are X and Z independent?
 - No
- Are they conditionally independent given Y?
 - Yes

- Common effect



X: Raining
Z: Ballgame
Y: Traffic

- Are X and Z independent?
 - Yes
- Are they conditionally independent given Y?
 - No



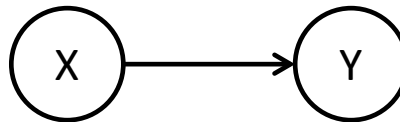
Compactness

- Suppose we have a Boolean variable X_i with k Boolean parents. How many rows does its conditional probability table have?
 - 2^k rows for all the combinations of parent values, each row requires one number p for $X_i = \text{true}$
- If each variable has no more than k parents, how many numbers does the complete network require?
 - $O(n \cdot 2^k)$ numbers – vs. $O(2^n)$ for the full joint distribution
- Example: How many nodes for the burglary network?
 $1 + 1 + 4 + 2 + 2 = 10$ numbers
(vs. specification of the complete joint probability $2^5 - 1 = 31$)

Constructing Bayesian networks

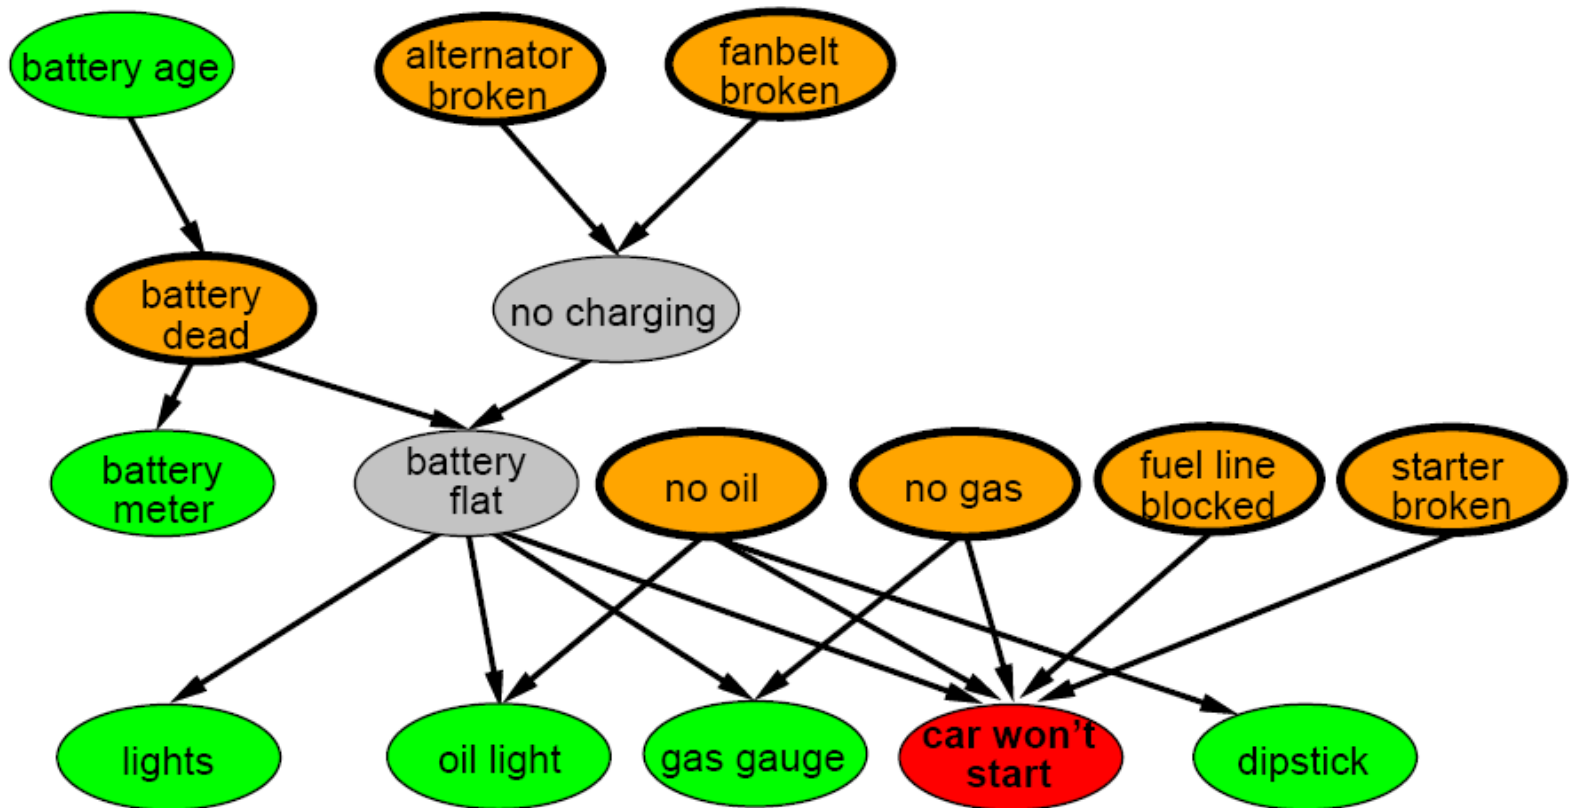
1. Choose an ordering of variables X_1, \dots, X_n
2. For $i = 1$ to n
 - add X_i to the network
 - select parents from X_1, \dots, X_{i-1} such that
$$P(X_i \mid \text{Parents}(X_i)) = P(X_i \mid X_1, \dots, X_{i-1})$$

Note: Networks are typically constructed by domain experts with causality in mind. E.g., X causes Y :



A more realistic Bayes Network: Car diagnosis

- **Initial observation:** car won't start
- **Green:** testable evidence
- **Orange:** "broken, so fix it" nodes
- **Gray:** "hidden variables" to ensure sparse structure, reduce parameters

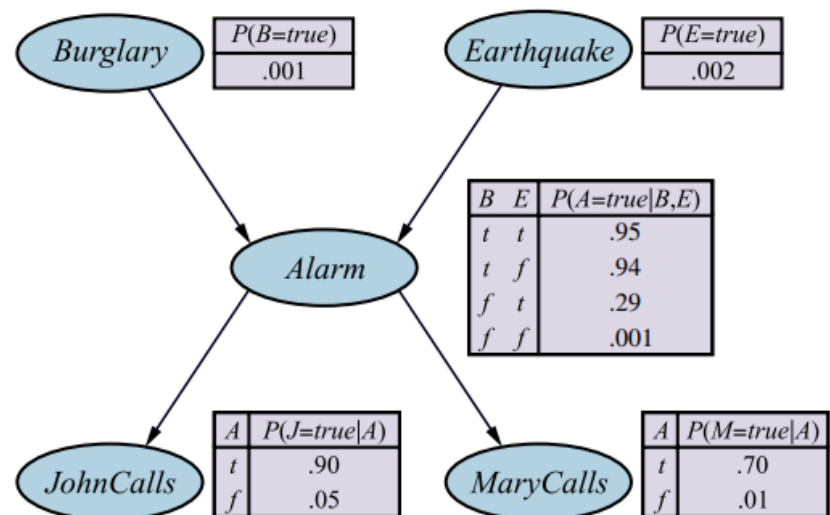


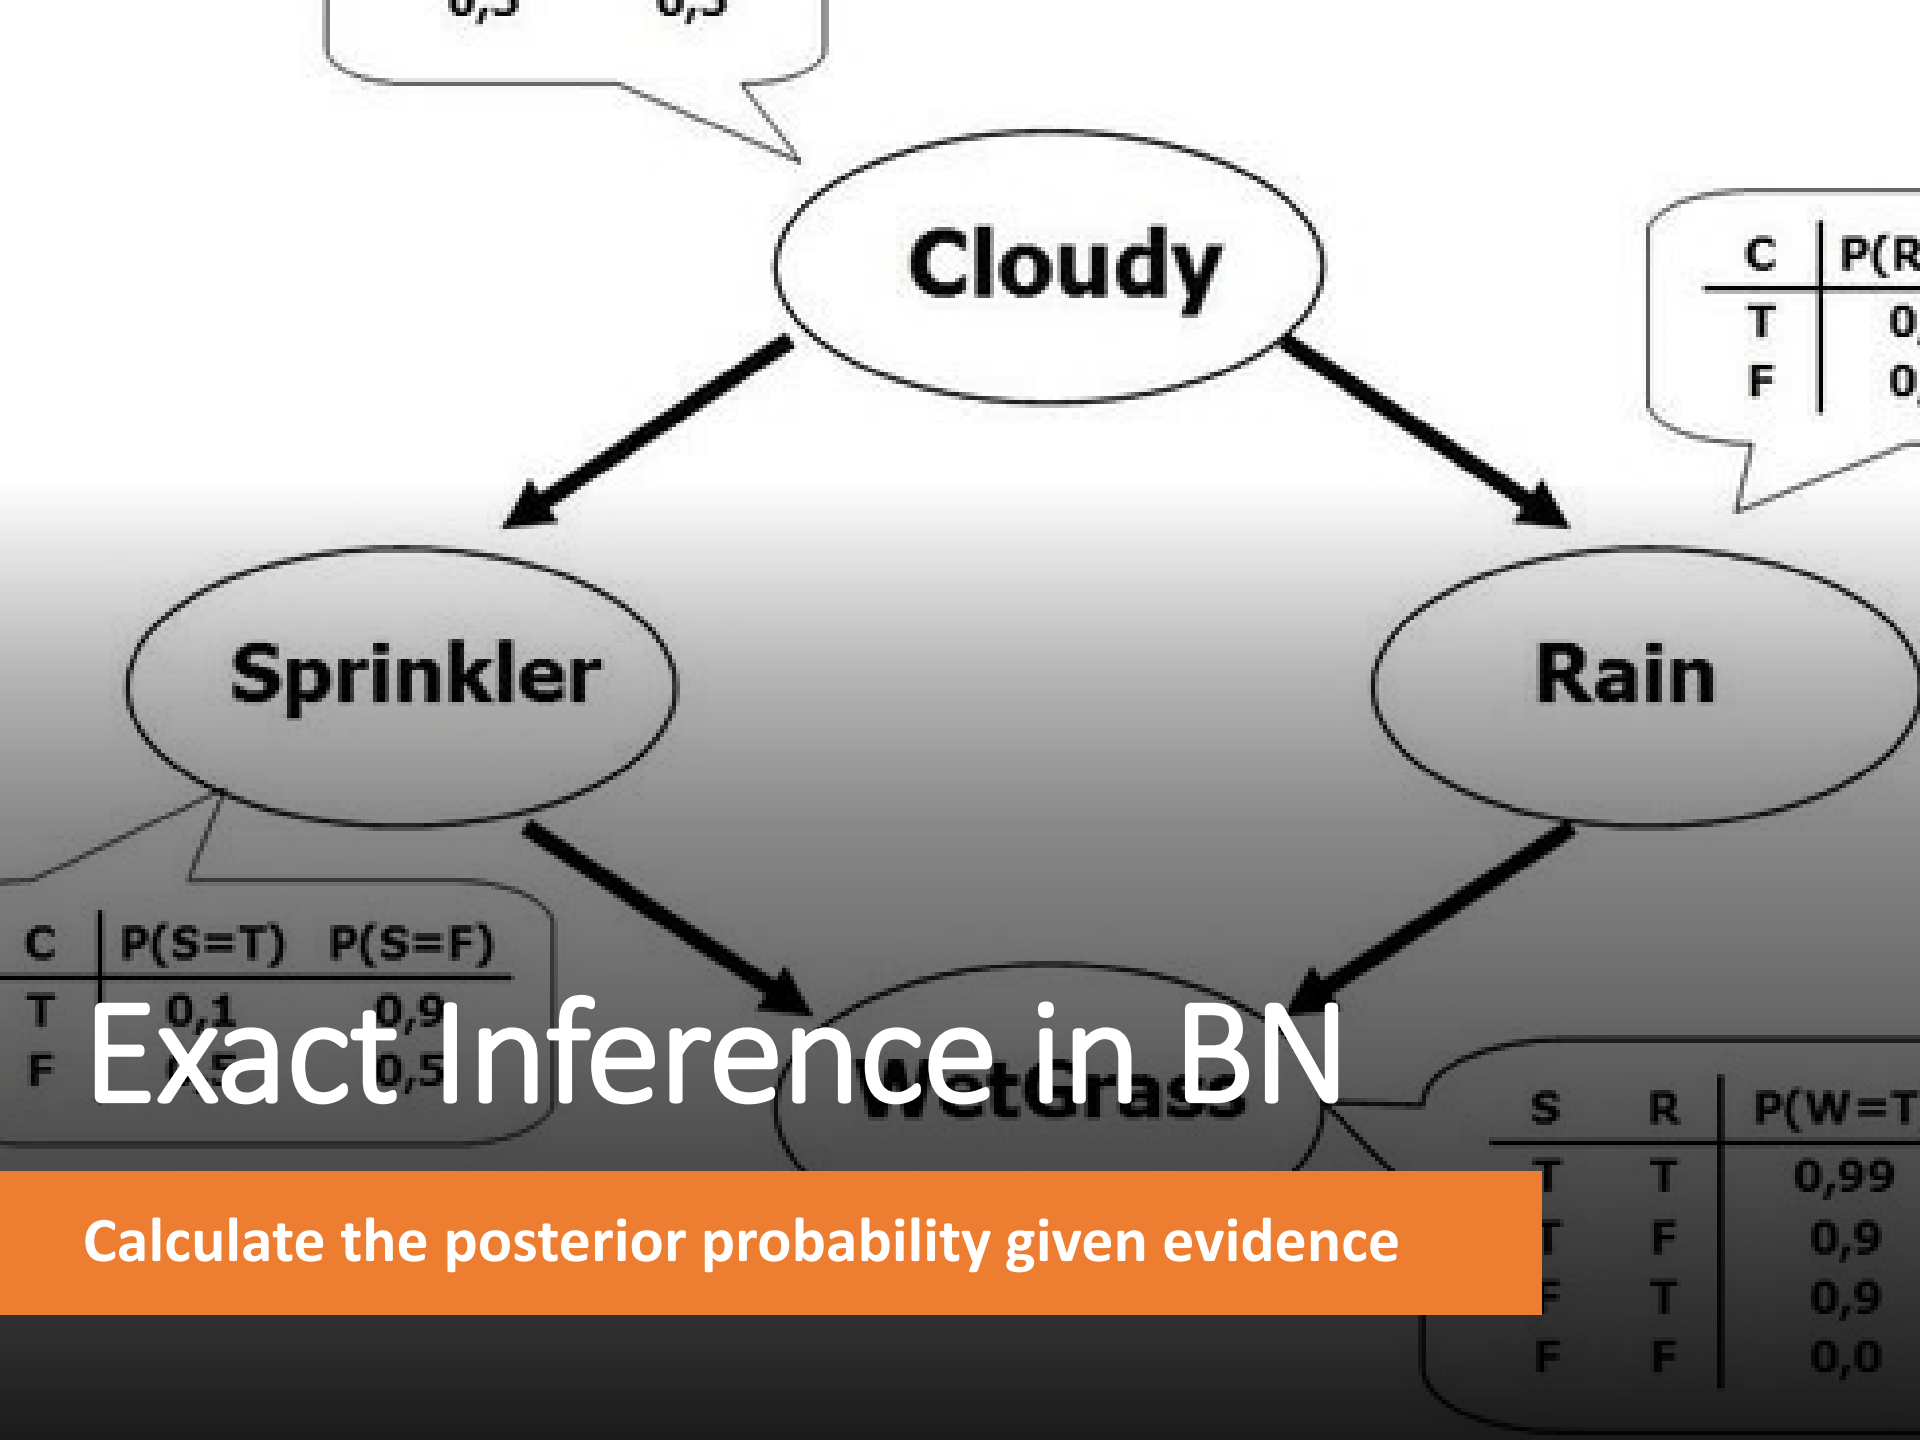
Summary

- Bayesian networks provide a natural representation for joint probabilities used to calculate conditional probabilities used in inference.
- Conditional independence (induced by causality) reduces the number of needed parameters.

$P(B, E, A, J, M)$ is defined by

- Representation
 - Topology
 - Conditional probability tables
 - Generally easy for domain experts to construct





Exact Inference in BN

Calculate the posterior probability given evidence

Exact Inference

Goal

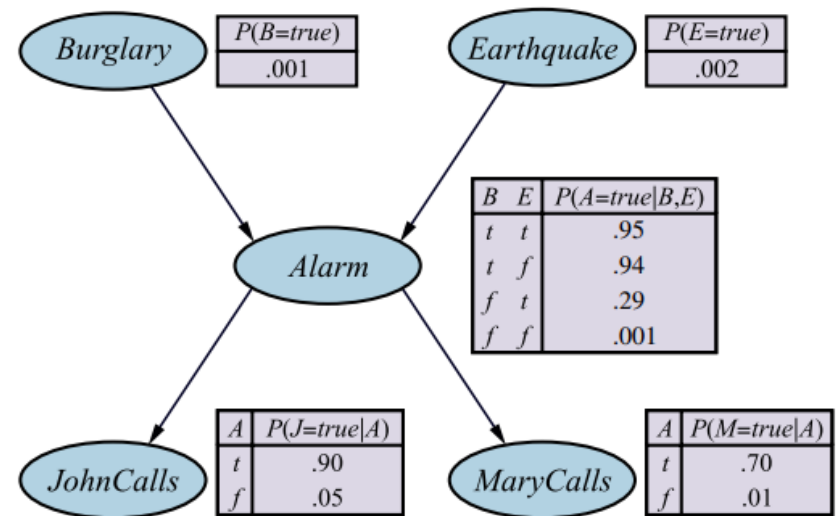
- Query variables: X
- Evidence (observed) variables: $E = e$
- Set of unobserved variables: Y
- Calculate the probability of X given e .

If we know the full joint distribution $P(X, E, Y)$, we can infer X by:

$$P(X|E = e) = \frac{P(X, e)}{P(e)} \propto \sum_y P(X, e, y)$$

Sum over values of unobservable variables = marginalizing them out.

Exact inference: Example



Assume we can observe being called and the two variables have the values j and m . We want to know the probability of a burglary.

Query: $P(B | j, m)$ with unobservable variables: Earthquake, Alarm

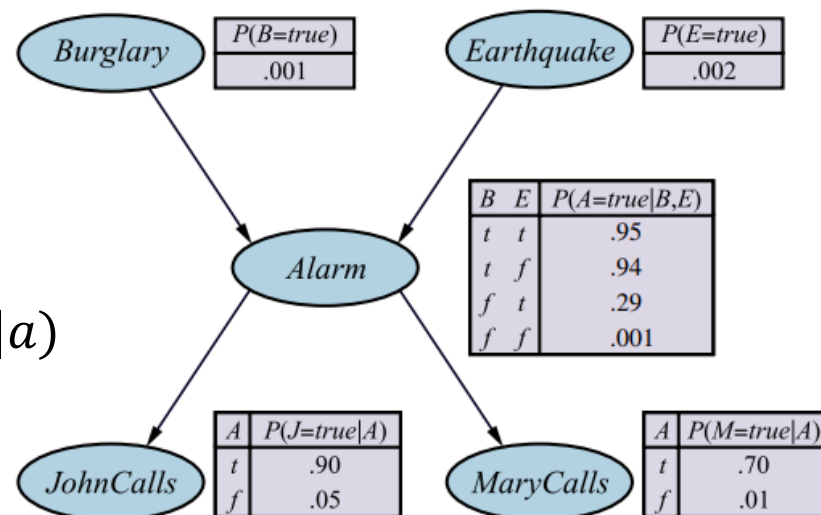
$$\begin{aligned}
 P(b|j, m) &= \frac{P(b, j, m)}{P(j, m)} \propto \sum_{E=e} \sum_{A=a} P(b, e, a, j, m) \\
 &= \sum_{E=e} \sum_{A=a} P(b)P(e)P(a|b, e)P(j|a)P(m|a) \\
 &= P(b) \sum_{E=e} P(e) \sum_{A=a} P(a|b, e) P(j|a)P(m|a)
 \end{aligned}$$

Full joint probability and marginalize over E and A

Exact inference: Example

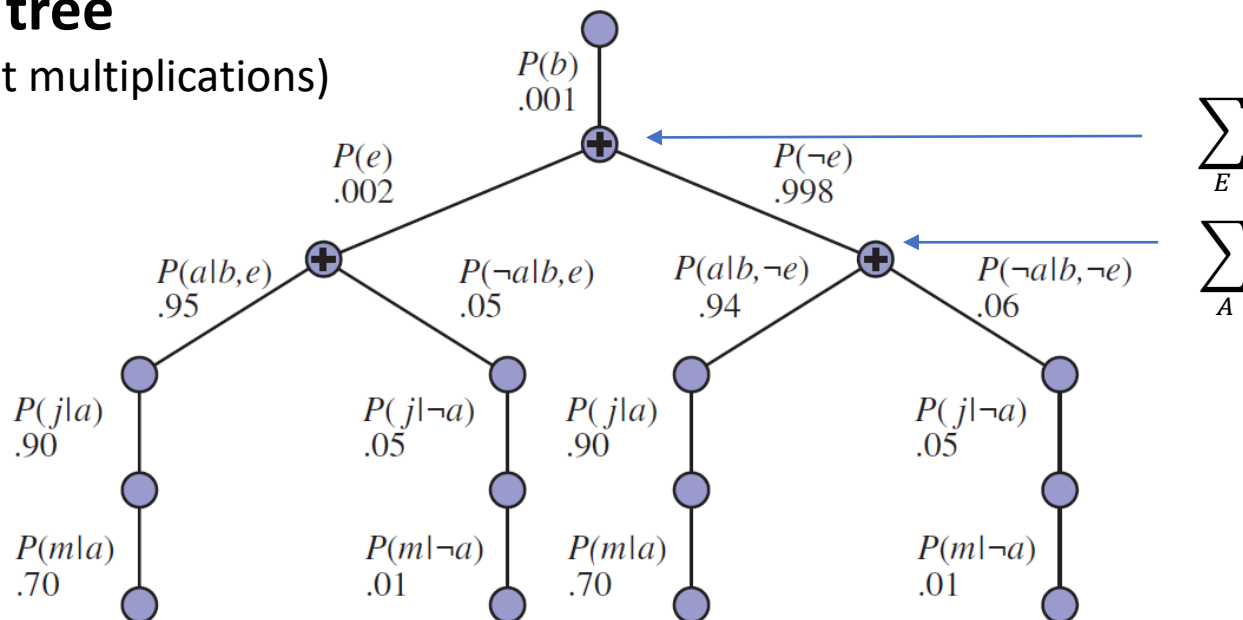
$$P(b|j, m)$$

$$\propto P(b) \sum_{E=e} P(e) \sum_{A=a} P(a|b, e) P(j|a) P(m|a)$$

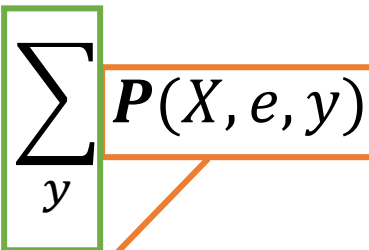


Evaluation tree

(lines represent multiplications)



Issues with Exact Inference in AI

$$P(X|E = e) = \frac{P(X, e)}{P(e)} \propto \sum_y P(X, e, y)$$


Problems


1. **Full joint distributions are too large** to store.

Bayes nets provide significant savings for representing the conditional probability structure.

2. Marginalizing out many unobservable variables Y may involve **too many summation terms**.

This summation is called **exact inference by enumeration**. Unfortunately, it does not scale well (#p-hard).

In praxis, **approximate inference by sampling** is used.




Approximate Inference in BN

Estimate the posterior probability given evidence

BN as a Generative Model



Bayesian networks can be used as ***generative models***.



Allows us to efficiently generate samples from the joint distribution.




Idea: Generate samples from the network to estimate joint and conditional probability distributions.

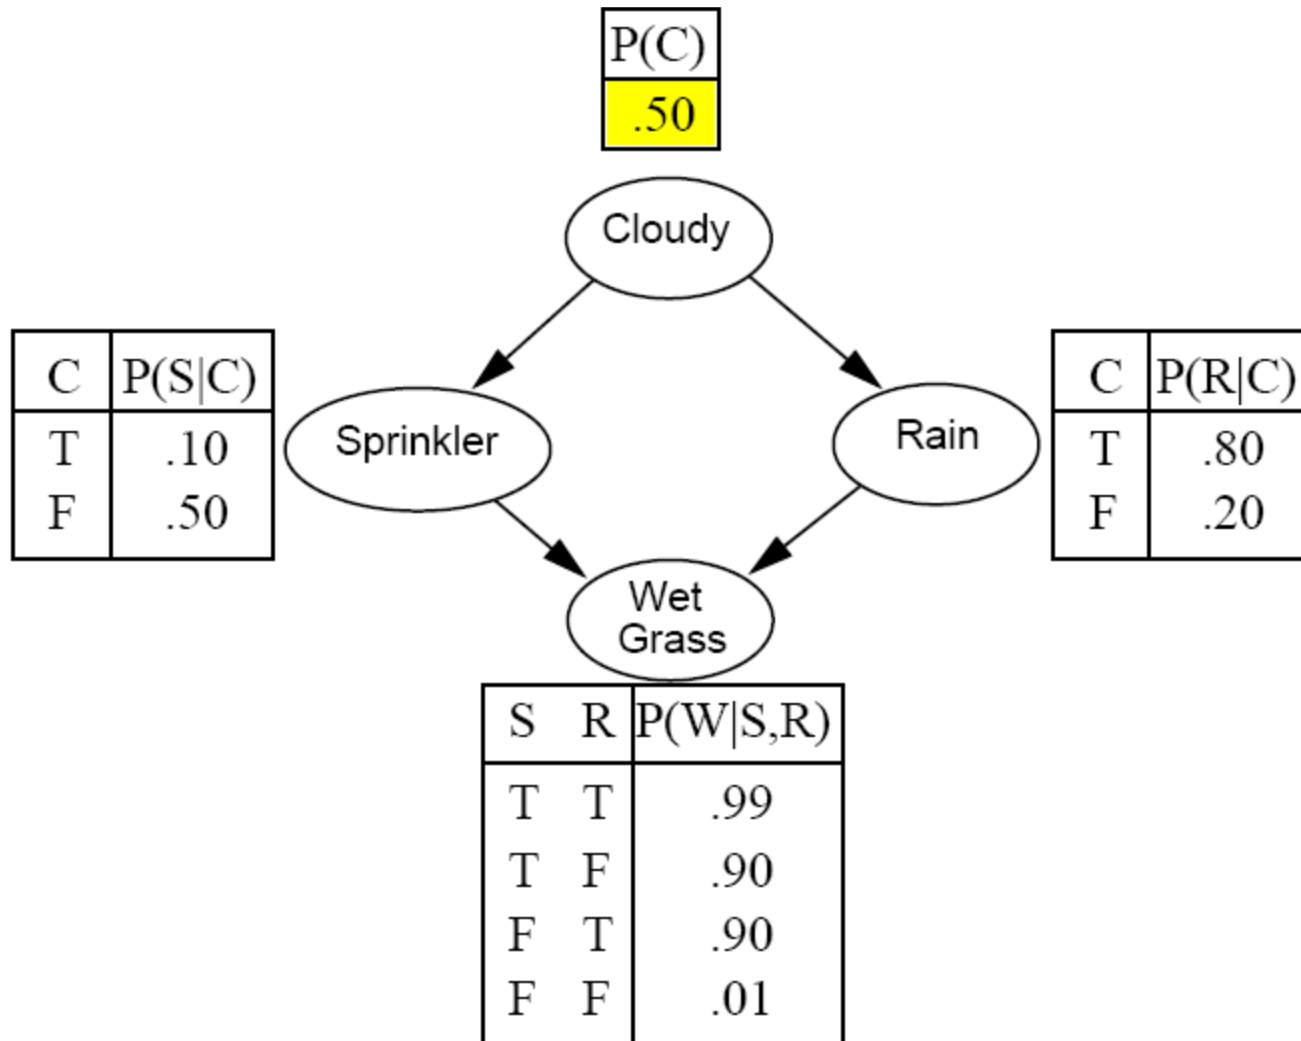
Prior-Sample Algorithm to Create a Sample (Event)

```
function PRIOR-SAMPLE( $bn$ ) returns an event sampled from the prior specified by  $bn$   
inputs:  $bn$ , a Bayesian network specifying joint distribution  $\mathbf{P}(X_1, \dots, X_n)$   
  
 $\mathbf{x} \leftarrow$  an event with  $n$  elements  
for each variable  $X_i$  in  $X_1, \dots, X_n$  do  
     $\mathbf{x}[i] \leftarrow$  a random sample from  $\mathbf{P}(X_i \mid \text{parents}(X_i))$   
return  $\mathbf{x}$ 
```

We need to start with the random variables that have no parents.

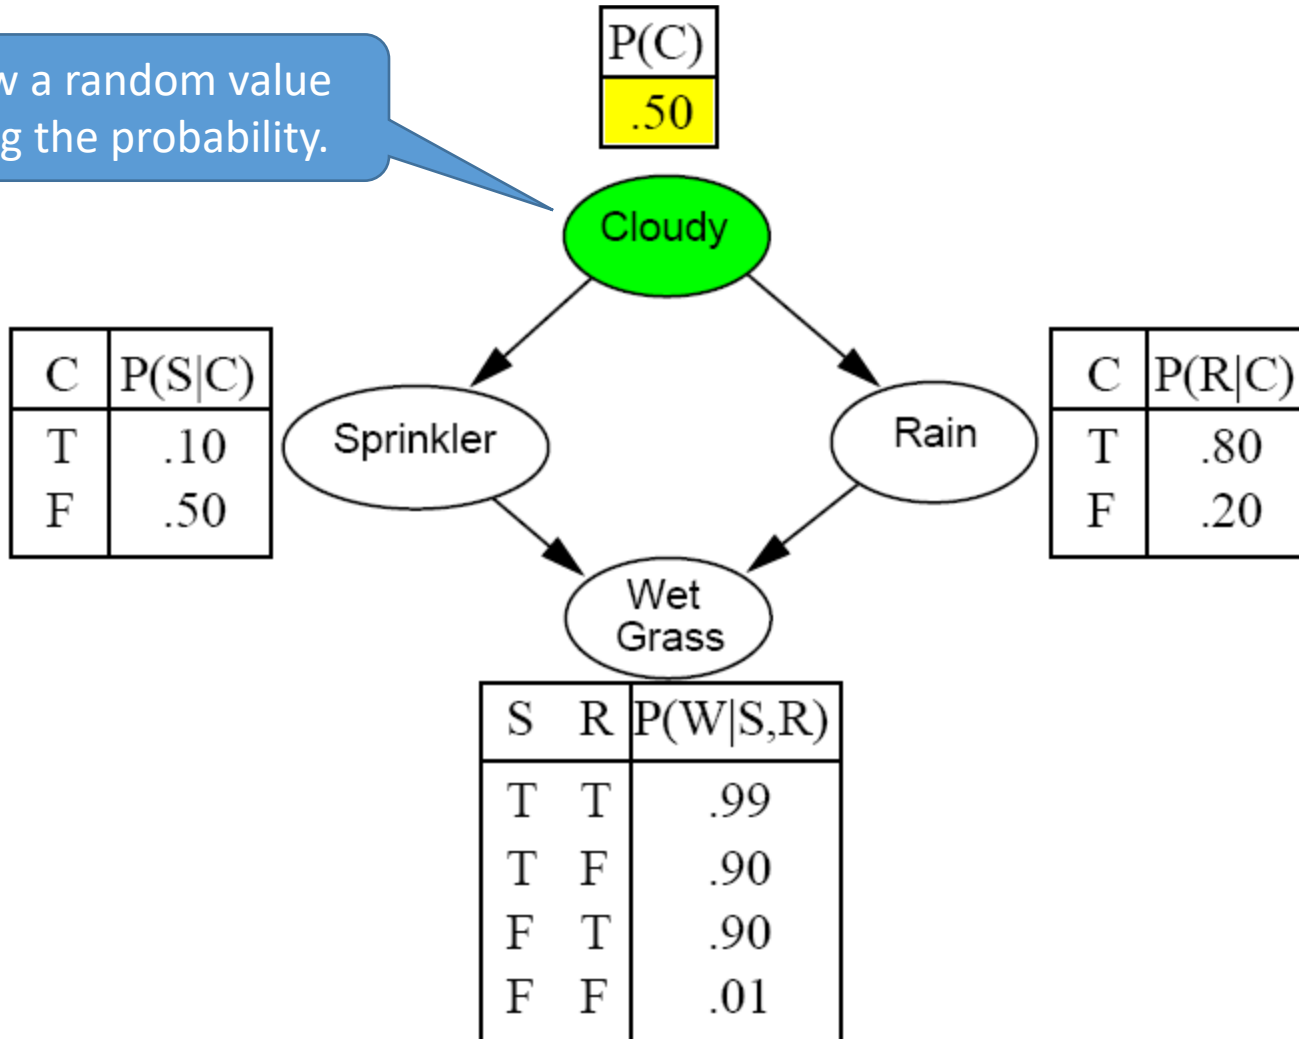


Example: Sampling from a Bayesian Network

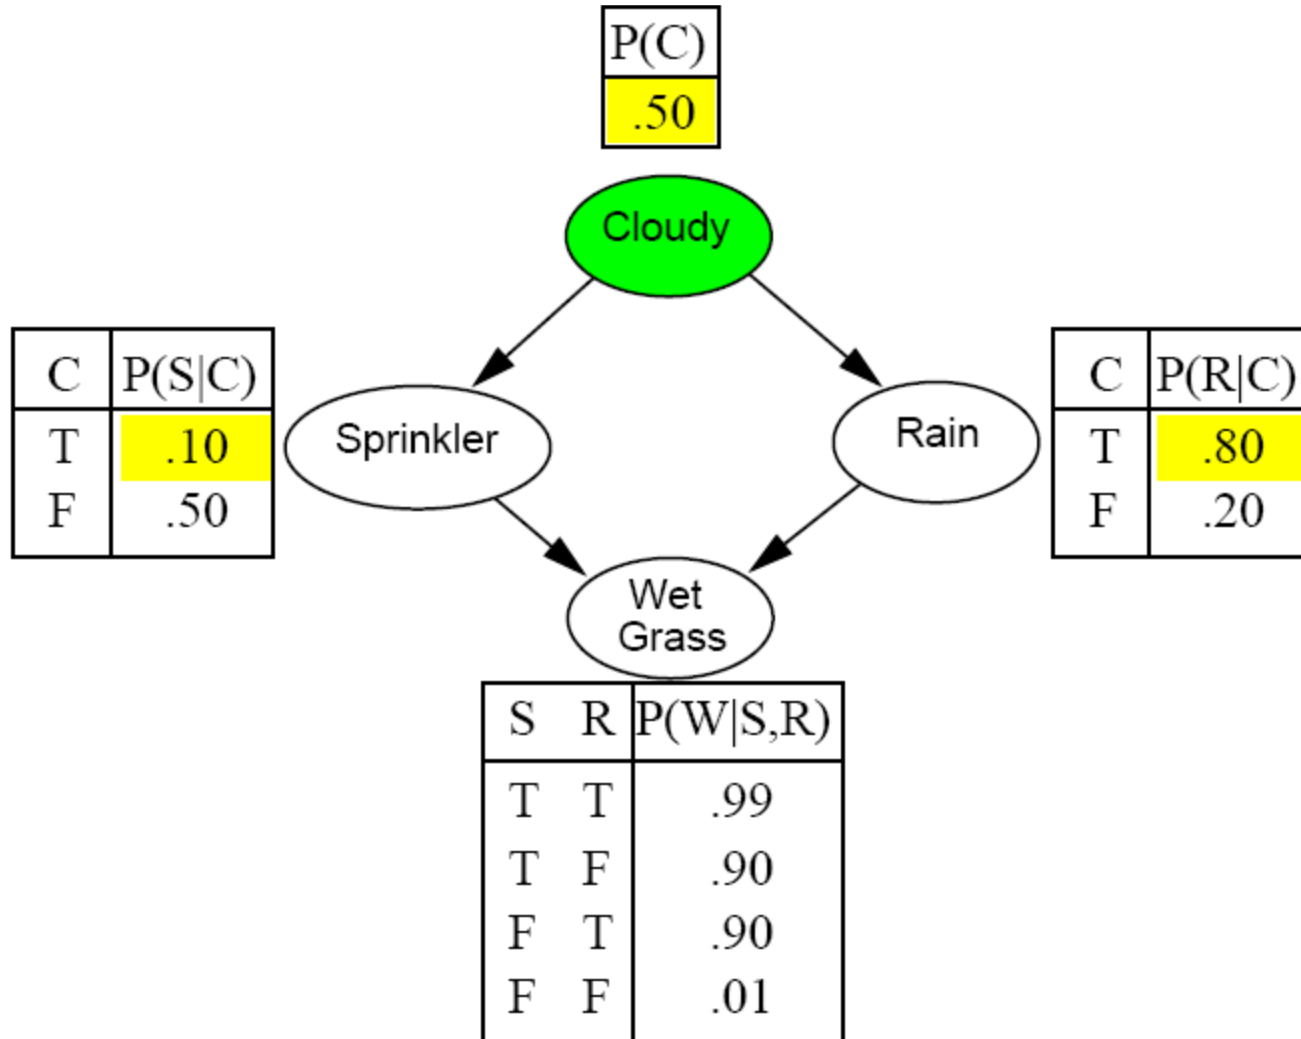


Example: Sampling from a Bayesian Network

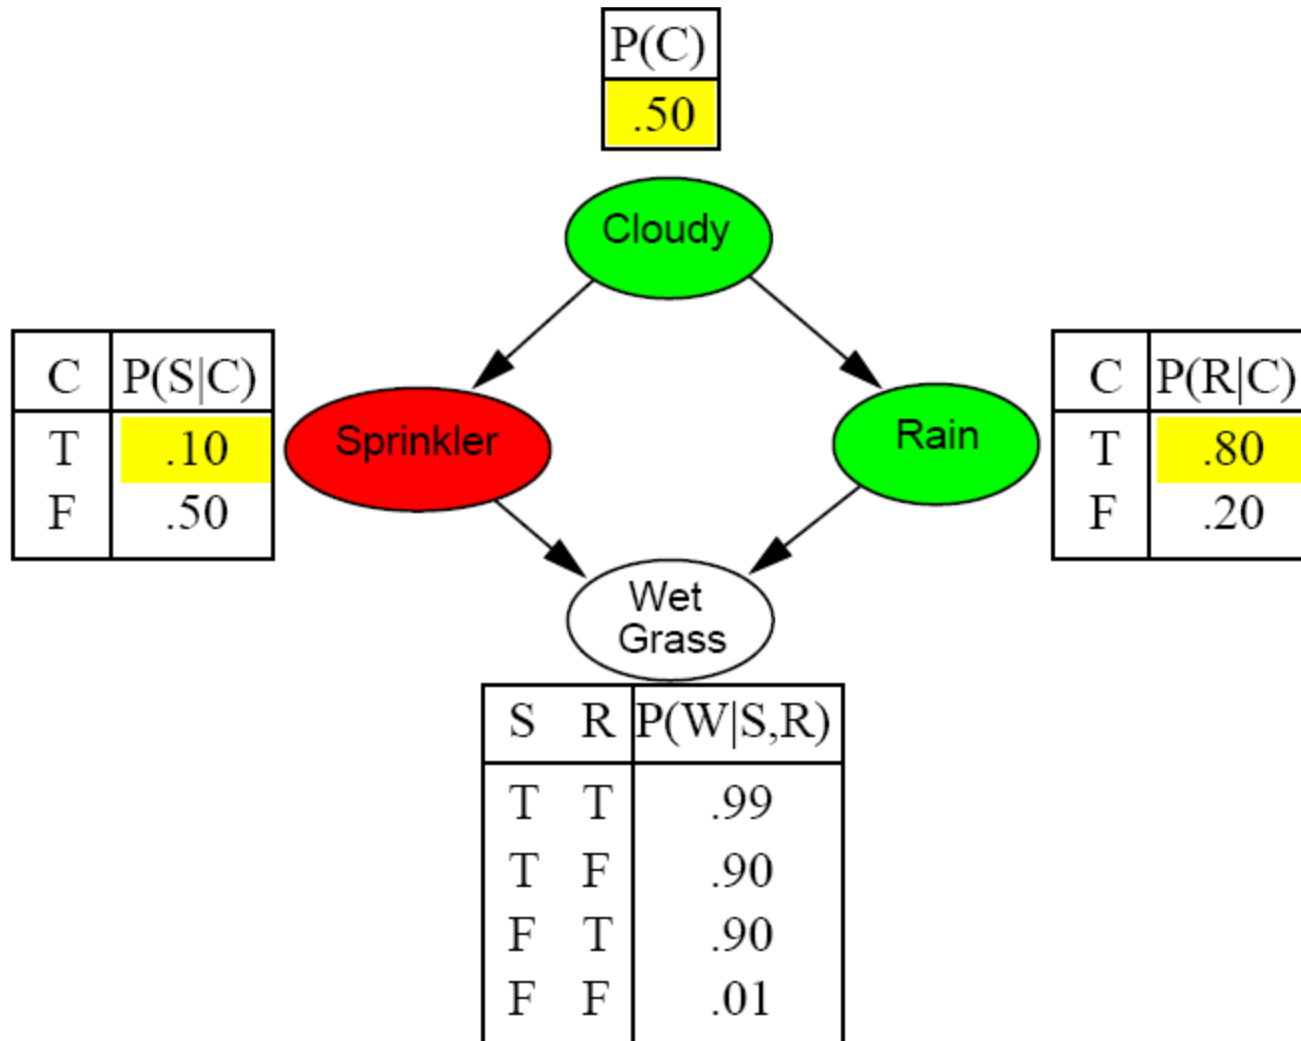
Draw a random value using the probability.



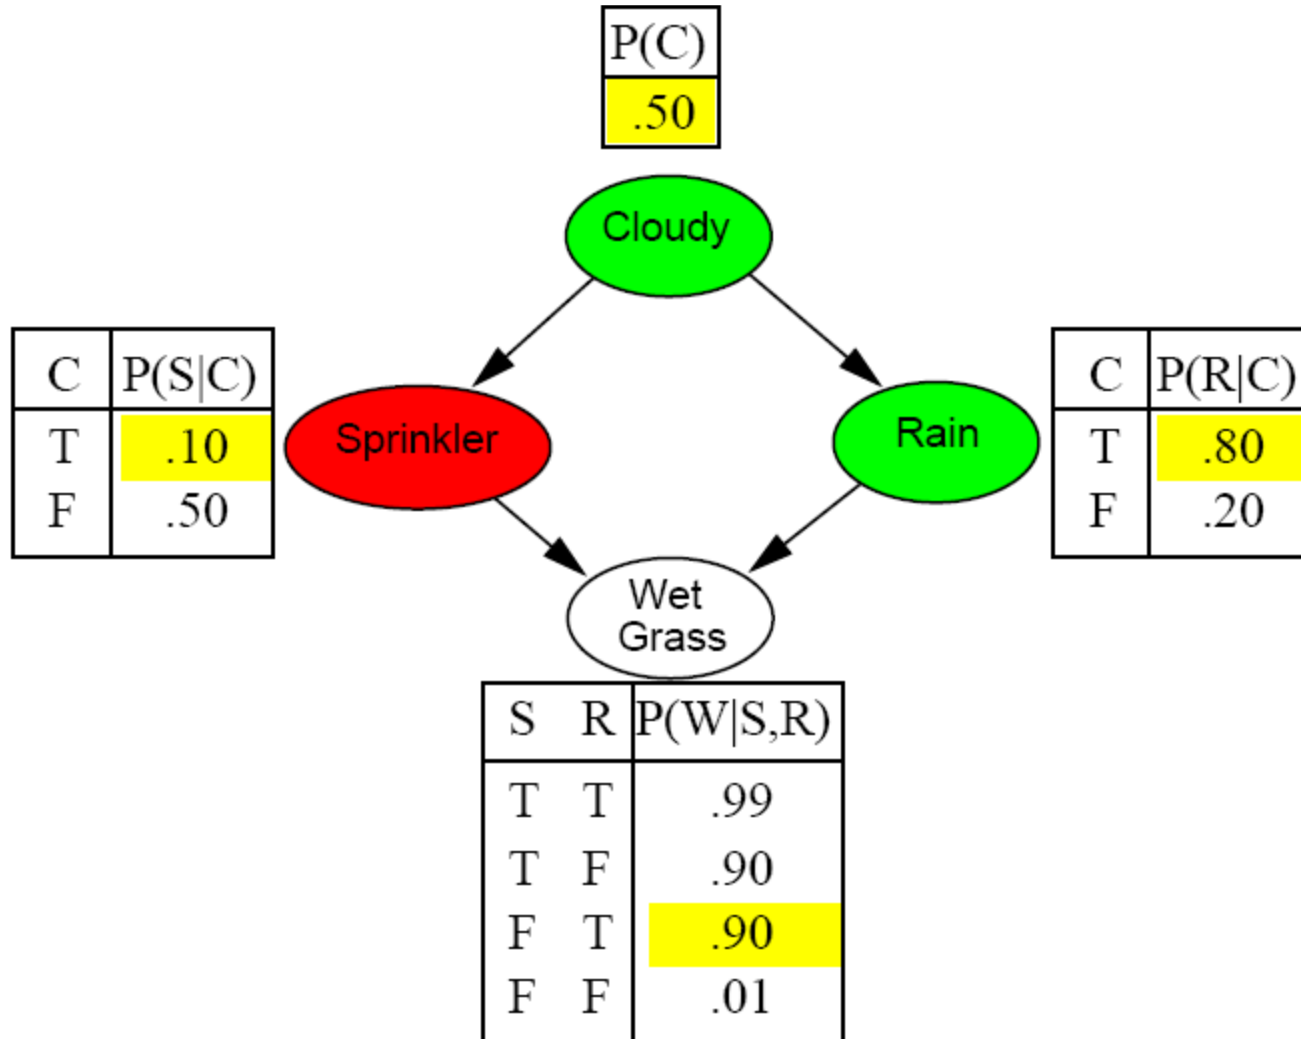
Example: Sampling from a Bayesian Network



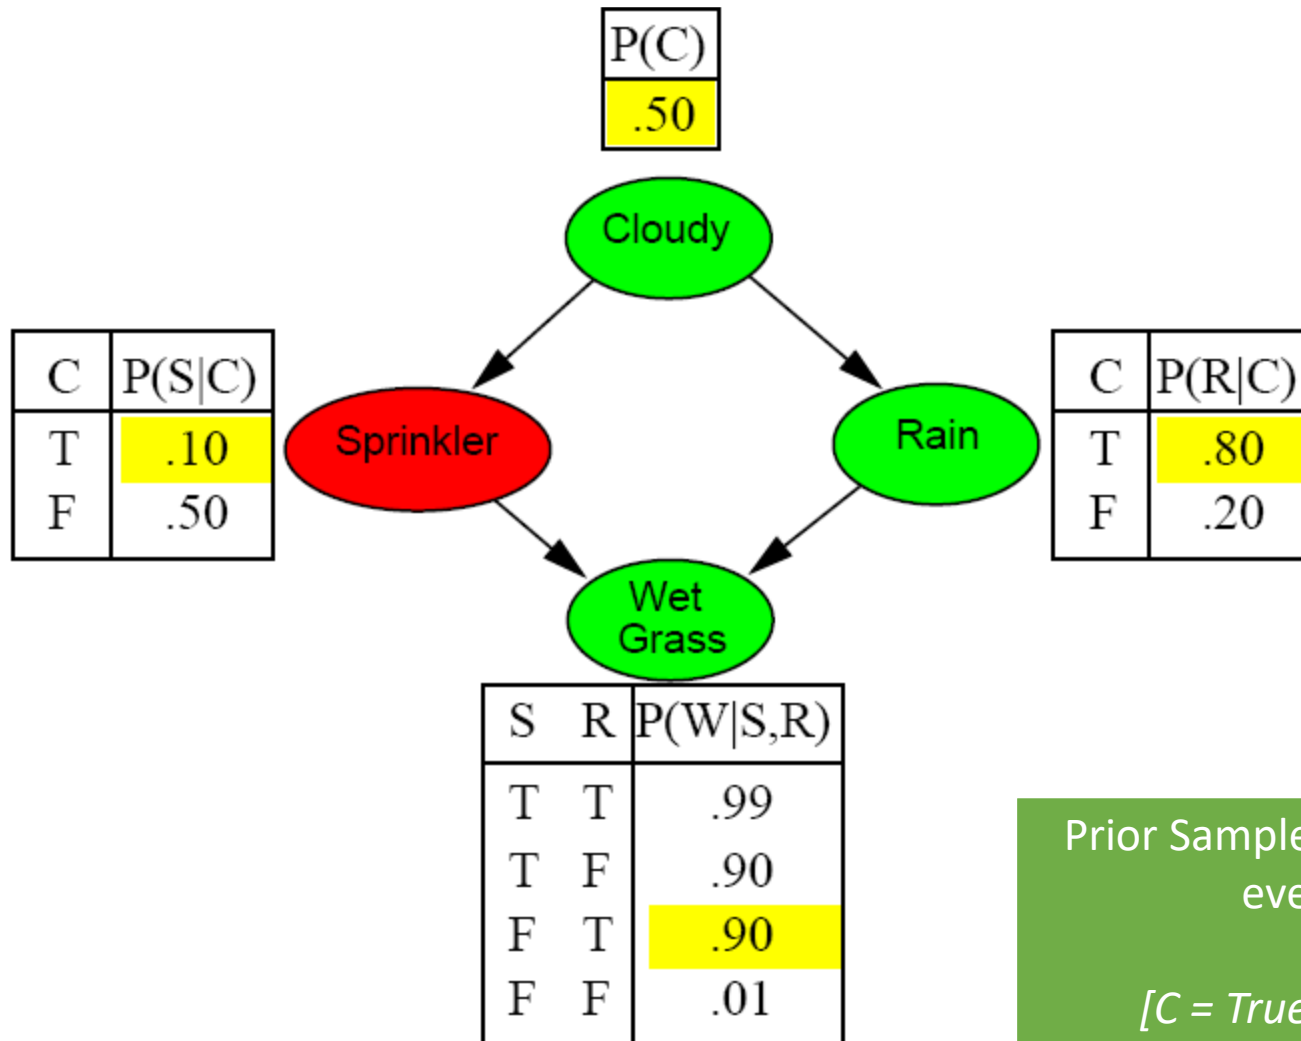
Example: Sampling from a Bayesian Network



Example: Sampling from a Bayesian Network



Example: Sampling from a Bayesian Network



Prior Sample returns the event:

$[C = \text{True}, S = \text{False}, R = \text{True}, W = \text{True}]$

Estimating the Joint Probability Distribution

Sample N times and determine $N_{PS}(x_1, x_2, \dots, x_n)$, the count of how many times Prior-Sample produces event (x_1, x_2, \dots, x_n) .

$$\hat{P}(x_1, x_2, \dots, x_n) = \frac{N_{PS}(x_1, x_2, \dots, x_n)}{N}$$

The marginal probability of partially specified event (some x values are known) can also be calculated. E.g.,

$$\hat{P}(x_1) = \frac{N_{PS}(x_1)}{N}$$

Estimating Conditional Probabilities:

Rejection sampling

Sample N times and **ignore the samples that are not consistent with the evidence e .**

$$\hat{P}(X|e) = \alpha N_{PS}(X, e) = \frac{N_{PS}(X, e)}{N_{PS}(e)}$$

Issue: What if e is a rare event?

- Example: burglary \wedge earthquake
- Rejection sampling ends up throwing away most of the samples. This is very inefficient!

Estimating Conditional Probabilities:

Rejection sampling

function REJECTION-SAMPLING(X, \mathbf{e}, bn, N) **returns** an estimate of $\mathbf{P}(X \mid \mathbf{e})$
inputs: X , the query variable
 \mathbf{e} , observed values for variables \mathbf{E}
 bn , a Bayesian network
 N , the total number of samples to be generated
local variables: \mathbf{C} , a vector of counts for each value of X , initially zero

for $j = 1$ **to** N **do**
 $\mathbf{x} \leftarrow \text{PRIOR-SAMPLE}(bn)$
 if \mathbf{x} is consistent with \mathbf{e} **then**
 $\mathbf{C}[j] \leftarrow \mathbf{C}[j] + 1$ where x_j is the value of X in \mathbf{x}
return NORMALIZE(\mathbf{C})

We throw away many samples
if \mathbf{e} is rare!

Estimating Conditional Probabilities: Importance sampling (likelihood weighting)

Goal: Avoid the need of rejection sampling to throw out samples.

1. Fix the evidence $E = e$ for sampling and estimate the probability for the non-evidence variables


$$Q_{WS}(x)$$

2. Correct the probabilities using weights
 $P(x|e) = w(x)Q_{WS}(x)$

Turns out the weights in this case can be easily calculated

$$w(x) = \alpha \prod_{i=1}^m P(e_i | \text{parents}(E_i))$$

Example: Evidence = it rains



Estimating Conditional Probabilities: Markov Chain Monte Carlo Sampling (MCMC)

- **Generates a sequence of samples** instead of creating each sample individually from scratch.
- Create a state by making random changes to the current state. The sequence of states forms a random process called a **Markov Chain** (MC).
- The MCs stationary distribution turns out to be the posterior distribution of the non-evidence variables.
- Estimate the stationary distribution using **Monte Carlo** simulation by counting how often each state is reached and normalize to obtain probability estimates.
- Algorithms:
 1. Gibbs sampling (works well for BNs)
 2. Metropolis-Hastings sampling

Note: Simulated annealing belongs to the family of MCMC algorithms.

Gibbs sampling in Bayes Networks

function GIBBS-ASK(X, \mathbf{e}, bn, N) **returns** an estimate of $\mathbf{P}(X | \mathbf{e})$

local variables: \mathbf{C} , a vector of counts for each value of X , initially zero

\mathbf{Z} , the nonevidence variables in bn

\mathbf{x} , the current state of the network, initialized from \mathbf{e}

initialize \mathbf{x} with random values for the variables in \mathbf{Z}

for $k = 1$ **to** N **do**

choose any variable Z_i from \mathbf{Z} according to any distribution $\rho(i)$

 set the value of Z_i in \mathbf{x} by sampling from $\mathbf{P}(Z_i | mb(Z_i))$

$\mathbf{C}[j] \leftarrow \mathbf{C}[j] + 1$ where x_j is the value of X in \mathbf{x}

return NORMALIZE(\mathbf{C})

Random
State

Change one
variable

Count

- $mb(Z_i)$ is the Markov blanket of random variable Z_i (all variables it can be dependent of, i.e., parents, children and parents of children).

$$P(z_i | mb(Z_i)) = \alpha P(z_i | \text{parents}(Z_i)) \prod_{Y_j \in \text{children}(X_i)} P(y_j | \text{parents}(Y_j))$$

Gibbs Sampling: Example

Find

$$P(\text{Rain} \mid \text{Sprinkler} = \text{true}, \text{WetGrass} = \text{true}).$$



Determine states and calculate transition probabilities of the Markov chain for the query using $P(z_i \mid mb(Z_i))$.

The algorithm wanders around in this graph using the stated transition probabilities.

Assume that we observe 20 states with *Rain* = *true* and 60 with *rain* = *false*:

$$\text{NORMALIZE}(\langle 20, 60 \rangle) = \langle 0.25, 0.75 \rangle$$

$$P(\text{Rain} \mid \text{Sprinkler} = \text{true}, \text{WetGrass} = \text{true}) \approx 0.75$$



Note the self-loops: the state stays the same when either variable is chosen and then resamples the same value it already has.



Conclusion

- Bayesian networks provide an efficient way to store a probabilistic model by exploiting (conditional) independence between variables.
- Exact Inference (estimating conditional probabilities) is difficult, for all but tiny models.
- State of the art is to use approximate inference by sampling from the model.
- Software libraries provide general inference engines.